



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

# Specifying and Checking File System Crash-Consistency Models

Steven Lang

September 4, 2016

# Motivation

The problem:

- ▶ POSIX file-system-interfaces do not define possible outcomes of a crash
- ▶ Can lead to
  - ▶ Corrupt application states
  - ▶ Catastrophic data loss

# Motivation


## Replace-Via-Rename Pattern

```
/* "file" has old data */  
fd = open("file.tmp");  
write(fd, new, size);  
close(fd);  
rename("file.tmp", "file");
```

# Motivation

## Replace-Via-Rename Pattern


```
/* "file" has old data */  
fd = open("file.tmp");  
write(fd, new, size);  
close(fd);  
rename("file.tmp", "file");
```



# Motivation

## Replace-Via-Rename Pattern

```
/* "file" has old data */
fd = open("file.tmp");
write(fd, new, size);
close(fd);
rename("file.tmp", "file");
```




file's on-disk state	possible executions seen on disk
new	open, write <sub>0</sub> , rename, write <sub>1</sub> , ...

# Motivation

## Replace-Via-Rename Pattern

```
/* "file" has old data */
fd = open("file.tmp");
write(fd, new, size);
close(fd);
rename("file.tmp", "file");
```




file's on-disk state	possible executions seen on disk
new	open, write <sub>0</sub> , rename, write <sub>1</sub> , ...
old	open, write <sub>0</sub> , crash

# Motivation

## Replace-Via-Rename Pattern

```
/* "file" has old data */  
fd = open("file.tmp");  
write(fd, new, size);  
close(fd);  
rename("file.tmp", "file");
```




file's on-disk state	possible executions seen on disk
new	open, write <sub>0</sub> , rename, write <sub>1</sub> , ...
old	open, write <sub>0</sub> , crash
empty	open, rename, crash

# Motivation

## Replace-Via-Rename Pattern

```
/* "file" has old data */
fd = open("file.tmp");
write(fd, new, size);
close(fd);
rename("file.tmp", "file");
```



file's on-disk state	possible executions seen on disk
new	open, write <sub>0</sub> , rename, write <sub>1</sub> , ...
old	open, write <sub>0</sub> , crash
empty	open, rename, crash
partial new	open, write <sub>0</sub> , rename, crash



Motivation

Introduction

Background

Crash-Consistency Models

FERRITE

Conclusion

Outlook

# Introduction

- ▶ Modern file system optimizations relax the order in which operations are executed

Good:

- ▶ Provide significant performance gains

Bad:

- ▶ Invisible to applications
- ▶ Machine crashes during out-of-order execution can harm the data's persistence

# Introduction

- ▶ Key challenge for application writers:
  - ▶ Understand the behavior of file systems across system crashes
- ▶ They make assumptions about crash guarantees provided by file systems
- ▶ They base their applications on these assumptions

# Introduction

- Being too optimistic about crash guarantees leads to serious **data losses**



# Introduction

- ▶ Being too optimistic about crash guarantees leads to serious **data losses**
- ▶ Being too conservative is **expensive** in energy, performance, and hardware lifespan



# Background

## The POSIX file system interface

- ▶ POSIX standard defines a set of system calls for fs access

operation	description
open	allocate file descriptor
write, read	perform file operations
link, unlink, mkdir	perform directory operations
sync, fsync	explicitly flush data to disk
close	deallocate file descriptor

- ▶ fsync system call is key to provide data integrity

# Crash-Consistency Models

- ▶ Used to define permissible states of a file system after a crash
- ▶ Consist of:
  - ▶ *Litmus tests*: demonstrate allowed/forbidden behaviors of file systems across crashes
  - ▶ *Formal specifications*: logic and state machines, describing crash-consistency behavior axiomatic and operational

# Crash-Consistency Models

## Litmus tests

Litmus tests consist of three parts:

1. Initial setup (optional)

Example:

initial:

```
f ← creat("f", 0600)
```



# Crash-Consistency Models

## Litmus tests

Litmus tests consist of three parts:

1. Initial setup (optional)
2. Main body

Example:

initial:

```
f ← creat("f", 0600)
```

main:

```
write(f, "data")
```

```
fsync(f)
```

```
mark("done")
```

```
close(f)
```

# Crash-Consistency Models

## Litmus tests

Litmus tests consist of three parts:

1. Initial setup (optional)
2. Main body
3. Final checking

Example:

initial:

```
f ← creat("f", 0600)
```

main:

```
write(f, "data")
```

```
fsync(f)
```

```
mark("done")
```

```
close(f)
```

exists?:

```
marked("done")
```

```
∧ content("f") ≠ "data"
```

# Crash-Consistency Models

## Litmus tests: Prefix-append (PA)

- ▶ The prefix-append (PA) litmus test checks whether, in the event of a crash, a file always contains a prefix of the data that has been appended to it

initial:

```
N ← 2500
as, bs ← "a" * N, "b" * N
f ← creat("file", 0600)
write(f, as)
```

main:

```
write(f, bs)
```

exists?:

```
content("file")  $\not\subseteq$  as + bs
```

- ▶ Also known as "safe-append"
- ▶ Popular file systems (ext4) do not guarantee this property

# Crash-Consistency Models

## Litmus tests: Atomic-replace-via-rename (ARVR)

- ▶ ARVR checks whether replacing file contents via rename is atomic across crashes

initial:

```
g ← creat("file", 0600)
write(g, old)
```

main:

```
f ← creat("file.tmp", 0600)
write(f, new)
rename("file.tmp", "file")
```

exists?:

```
content("file") ≠ old
∧ content("file") ≠ new
```

# Crash-Consistency Models

## Formal specifications

Two styles of specification:

- ▶ *axiomatic*: describe valid crash behaviors declaratively, using a set of axioms and ordering relations
- ▶ *operational*: abstract machines that simulate relevant aspects of file system behavior

# Crash-Consistency Models

## Formal specifications

Two styles of specification:

- ▶ *axiomatic*: describe valid crash behaviors declaratively, using a set of axioms and ordering relations
- ▶ *operational*: abstract machines that simulate relevant aspects of file system behavior

Example models will be shown for:

- ▶ seqfs, an ideal file system with strong crash consistency guarantees
- ▶ ext4, a real file system with weak consistency guarantees

# Crash-Consistency Models

## Axiomatic specifications

- ▶ Consist of a set of rules/axioms, that specify whether a given execution of a program is allowed

## Definition (Program)

- ▶ Sequence of atomic events, updating an underlying file system

# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$



# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$
- ▶  $\sigma_{meta}$ : map from (*object identifiers*)  $\mapsto \mathbb{N}$

# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$
- ▶  $\sigma_{meta}$ : map from (*object identifiers*)  $\mapsto \mathbb{N}$
- ▶  $\sigma_{data}$ : map from  $\mathbb{N} \mapsto$  (*file system objects*)

# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$
- ▶  $\sigma_{meta}$ : map from (*object identifiers*)  $\mapsto \mathbb{N}$
- ▶  $\sigma_{data}$ : map from  $\mathbb{N} \mapsto$  (*file system objects*)
- ▶ Object identifiers  $i \in \mathbb{I}$

# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$
- ▶  $\sigma_{meta}$ : map from (*object identifiers*)  $\mapsto \mathbb{N}$
- ▶  $\sigma_{data}$ : map from  $\mathbb{N} \mapsto$  (*file system objects*)
- ▶ Object identifiers  $i \in \mathbb{I}$
- ▶ Object:  $\sigma(i) = \sigma_{data}(\sigma_{meta}(i))$

# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$
- ▶  $\sigma_{meta}$ : map from (*object identifiers*)  $\mapsto \mathbb{N}$
- ▶  $\sigma_{data}$ : map from  $\mathbb{N} \mapsto$  (*file system objects*)
- ▶ Object identifiers  $i \in \mathbb{I}$
- ▶ Object:  $\sigma(i) = \sigma_{data}(\sigma_{meta}(i))$
- ▶  $\sigma(i) = \perp$  if object  $\sigma(i)$  has not been created

# Axiomatic specifications

## File Systems

### Definition

- ▶ *file system*  $\sigma = \langle \sigma_{meta}, \sigma_{data} \rangle$
- ▶  $\sigma_{meta}$ : map from (*object identifiers*)  $\mapsto \mathbb{N}$
- ▶  $\sigma_{data}$ : map from  $\mathbb{N} \mapsto$  (*file system objects*)
- ▶ Object identifiers  $i \in \mathbb{I}$
- ▶ Object:  $\sigma(i) = \sigma_{data}(\sigma_{meta}(i))$
- ▶  $\sigma(i) = \perp$  if object  $\sigma(i)$  has not been created
- ▶ File object:  $\sigma(f) = \langle b, m \rangle$ 
  - $b$ : finite string of bits
  - $m$ : finite key-value map of file metadata

# Axiomatic specifications

## Events

### Definition

- ▶ Atomic access to a file system  $\sigma$
- ▶ *Update events*: modify objects in  $\sigma$
- ▶ *Synchronization events*: synchronize accesses to (parts of)  $\sigma$

# Axiomatic specifications

## Events: Update

Update events include writes to file (meta-) data:

event	operation
$write(f, a, d)$	$b[a] = d$ for $\sigma(f) = \langle b, m \rangle$
$setattr(f, k, v)$	$m[k] = v$ for $\sigma(f) = \langle b, m \rangle$
$extend(f, a, d, s)$	$m["size"] = s$ and $b[a] = d$
$link(i_1, i_2)$	$\sigma_{meta}(i_2) = \sigma_{meta}(i_1)$
$unlink(i)$	$\sigma_{meta}(i) = \perp$



# Axiomatic specifications

## Events: Synchronization

**Synchronization events** include write barriers, non-file-system events (e.g. network-messages) and begin/end of a transaction:

event	operation
<i>fsync</i> ( <i>i</i> )	sync. accesses to $\sigma(i)$
<i>sync</i> ()	<i>fsync</i> ( <i>i</i> ), $\forall i \in \mathbb{I}$
<i>mark</i> ( <i>l</i> )	mark event identified by unique label <i>l</i>
<i>begin</i> ()	begins a new transaction
<i>commit</i> ()	ends the current transaction

# Axiomatic specifications

## Traces

### Definition

- ▶  $t_P$ : sequence of file system events of a program  $P$
- ▶  $\leq_{t_P}$ : total order on events induced by the trace  $t_P$  :  $e_1 \leq_{t_P} e_2$  iff  $e_1$  occurs before  $e_2$  in the trace  $t_P$
- ▶  $\tau_P$ : *canonical trace* of  $P$ , which is free of crashes and is a strict sequential execution of  $P$
- ▶  $c_P$ : *crash trace* - prefix of a *valid* trace  $t_P$  that respects transactional semantics:  
 $c_P$  contains the same number of *begin* and *commit* events

# Axiomatic specifications

## Traces

Trace  $t_P$  is valid, iff:

- ▶  $t_P$  is a permutation of  $\tau_P$
- ▶  $t_P$  respects sync. semantics of  $\tau_P$ :  
 $e_i \leq_{t_P} e_j$  when  $e_i \leq_{\tau_P} e_j$  and any of the following hold:
  - ▶  $e_i$  is an *fsync*, *sync*, *mark*, *begin* or *commit* event
  - ▶  $e_j$  is a *sync*, *mark*, *begin*, or *commit* event
  - ▶  $e_j$  is an *fsync* event on  $i$  and  $e_i$  is an update event on  $i$
- ▶  $t_P$  respects the update semantics of  $\tau_P$ :  
state of  $\sigma$  after applying  $t_P$  = state of  $\sigma$  after applying  $\tau_P$

# Axiomatic specifications

## Example: Ordered-two-file-overwrites

initial:

```
f ← creat("f", 0600)
g ← creat("g", 0600)
write(f, "0")
write(g, "0")
```

main:

```
pwrite(f, "1", 0)
pwrite(g, "1", 0)
fsync(g)
```

exists?:

```
content("f") = "0"
∧ content("g") = "1"
```

► initial state:

- $\sigma(f) = \sigma(g) = \langle b, m \rangle$
- $b = "0"$
- $m = \{ \text{permission} \mapsto "0600", \dots \}$

►  $\tau_P = [e_0, e_1, e_2]$ , with:

- $e_0 = \text{write}(f, 0, "1")$
- $e_1 = \text{write}(g, 0, "1")$
- $e_2 = \text{fsync}(g)$

► valid traces:

- $t_0 = [e_1, e_0, e_2]$
- $t_1 = [e_1, e_2, e_0]$

# Crash-Consistency Models

## Definition

- ▶ Crash-consistency models determine which valid program traces are permissible
- ▶ Strongest model:  
*Sequential Crash-Consistency* (SCC)
  - ▶ Permits no re-ordering of events
  - ▶ Only valid trace:  $\tau_P$

# Crash-Consistency Models

## Definition

- ▶ Crash-consistency models determine which valid program traces are permissible
- ▶ Strongest model:  
*Sequential Crash-Consistency* (SCC)
  - ▶ Permits no re-ordering of events
  - ▶ Only valid trace:  $\tau_P$
- ▶ **But:** Real file systems implement weaker models with additional valid traces

# Crash-Consistency Models

## Definition

- ▶ Crash-consistency models determine which valid program traces are permissible
- ▶ Strongest model:  
*Sequential Crash-Consistency (SCC)*
  - ▶ Permits no re-ordering of events
  - ▶ Only valid trace:  $\tau_P$
- ▶ **But:** Real file systems implement weaker models with additional valid traces

## Definition

- ▶ Let  $M$  be a crash-consistency model
- ▶  $M$  permits  $t_P \Leftrightarrow M(t_P, \tau_P) = \text{TRUE}$

# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

### Definition

- ▶  $SCC(t_P, \tau_P) = TRUE \Leftrightarrow t_P = \tau_P$



# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

### Definition

- $SCC(t_P, \tau_P) = TRUE \Leftrightarrow t_P = \tau_P$

Example: Ordered-two-file-overwrites

initial:

```
f ← creat("f", 0600)
g ← creat("g", 0600)
write(f, "0")
write(g, "0")
```

main:

```
pwrite(f, "1", 0)
pwrite(g, "1", 0)
fsync(g)
```

exists?:

```
content("f") = "0"
∧ content("g") = "1"
```

# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

Check whether SCC allows the surprising behavior of this litmus test:

1. Build all valid traces: SCC only permits

$$\tau_P = [write_f, write_g, fsync_g]$$

2. Build all crash-traces (prefixes) of the valid traces:

# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

Check whether SCC allows the surprising behavior of this litmus test:

1. Build all valid traces: SCC only permits

$$\tau_P = [write_f, write_g, fsync_g]$$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$c_{P_0}$	<i>crash</i>	"0"	"0"

# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

Check whether SCC allows the surprising behavior of this litmus test:

1. Build all valid traces: SCC only permits

$$\tau_P = [write_f, write_g, fsync_g]$$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$c_{P_0}$	<i>crash</i>	"0"	"0"
$c_{P_1}$	<i>write<sub>f</sub>, crash</i>	"1"	"0"

# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

Check whether SCC allows the surprising behavior of this litmus test:

1. Build all valid traces: SCC only permits

$$\tau_P = [write_f, write_g, fsync_g]$$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$c_{P_0}$	<i>crash</i>	"0"	"0"
$c_{P_1}$	<i>write<sub>f</sub>, crash</i>	"1"	"0"
$c_{P_2}$	<i>write<sub>f</sub>, write<sub>g</sub>, crash</i>	"1"	"1"

# Crash-Consistency Models

## Sequential Crash-Consistency (SCC)

Check whether SCC allows the surprising behavior of this litmus test:

1. Build all valid traces: SCC only permits

$$\tau_P = [write_f, write_g, fsync_g]$$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$c_{P_0}$	<i>crash</i>	"0"	"0"
$c_{P_1}$	<i>write<sub>f</sub>, crash</i>	"1"	"0"
$c_{P_2}$	<i>write<sub>f</sub>, write<sub>g</sub>, crash</i>	"1"	"1"

$\Rightarrow$  SCC does not allow

$$\text{content}("f") = "0" \wedge \text{content}("g") = "1"$$

# Crash-Consistency Models

## Definition

### Definition (Weaker Crash-Consistency Model)

- ▶ Model  $M_1$  is weaker than  $M_2$  iff  $M_1$  permits a superset of the valid traces permitted by  $M_2$
- ▶  $M_2(t_P, \tau_P) \Rightarrow M_1(t_P, \tau_P)$

# Crash-Consistency Models

## ext4 Crash-Consistency

### Definition

- ▶  $t_P$  is ext4 *crash-consistent* iff  $e_i \leq_{t_P} e_j$ ,  $\forall e_i, e_j$  such that  $e_i \leq_{\tau_P} e_j$  and at least one of the following conditions holds:
  1.  $e_i$  and  $e_j$  are metadata updates to the same file
  2.  $e_i$  and  $e_j$  are writes to the same block in the same file
  3.  $e_i$  and  $e_j$  are updates to the same directory
  4.  $e_i$  is a write and  $e_j$  is an extend to the same file



# Crash-Consistency Models

## ext4 Crash-Consistency example: Ordered-two-file-overwrites

Check whether ext4 allows the surprising behavior of this litmus test:

1. Build all valid traces - ext4 permits:

- ▶  $\tau = [write_f, write_g, fsync_g]$
- ▶  $t_0 = [write_g, write_f, fsync_g]$
- ▶  $t_1 = [write_g, fsync_g, write_f]$

# Crash-Consistency Models

## ext4 Crash-Consistency example: Ordered-two-file-overwrites

Check whether ext4 allows the surprising behavior of this litmus test:

1. Build all valid traces - ext4 permits:

- ▶  $\tau = [write_f, write_g, fsync_g]$
- ▶  $t_0 = [write_g, write_f, fsync_g]$
- ▶  $t_1 = [write_g, fsync_g, write_f]$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$c_{P_0}$	<i>crash</i>	"0"	"0"

# Crash-Consistency Models

## ext4 Crash-Consistency example: Ordered-two-file-overwrites

Check whether ext4 allows the surprising behavior of this litmus test:

1. Build all valid traces - ext4 permits:

- ▶  $\tau = [write_f, write_g, fsync_g]$
- ▶  $t_0 = [write_g, write_f, fsync_g]$
- ▶  $t_1 = [write_g, fsync_g, write_f]$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$CP_0$	<i>crash</i>	"0"	"0"
$CP_1$	<i>write<sub>g</sub>, crash</i>	"0"	"1"

# Crash-Consistency Models

## ext4 Crash-Consistency example: Ordered-two-file-overwrites

Check whether ext4 allows the surprising behavior of this litmus test:

1. Build all valid traces - ext4 permits:

- ▶  $\tau = [write_f, write_g, fsync_g]$
- ▶  $t_0 = [write_g, write_f, fsync_g]$
- ▶  $t_1 = [write_g, fsync_g, write_f]$

2. Build all crash-traces (prefixes) of the valid traces:

	crashtrace	content(f)	content(g)
$CP_0$	<i>crash</i>	"0"	"0"
$CP_1$	<i>write<sub>g</sub>, crash</i>	"0"	"1"

$\Rightarrow$  ext4 allows

$content("f") = "0" \wedge content("g") = "1"$

# Crash-Consistency Models

## Operational specifications

- ▶ Non-deterministic state machine  $M$
- ▶  $M$  is modeled as  $\langle \sigma, p \rangle$ , with  $\sigma$  as the file system and  $p$  as a program counter for  $P$

# Crash-Consistency Models

## Operational specifications: Examples

$$\frac{\sigma' = \text{FLUSH}(\text{APPLY}(P[p], \sigma))}{\langle \sigma, p \rangle \mapsto \langle \sigma', p + 1 \rangle} \text{STEPSEQ}$$

$$\frac{}{\langle \sigma, p \rangle \mapsto \langle \sigma, \perp \rangle} \text{CRASH}$$

**Figure:** An operational model for SCC

# Crash-Consistency Models

## Operational specifications: Examples

$$\frac{\sigma' = \text{FLUSH}(\text{APPLY}(P[p], \sigma))}{\langle \sigma, p \rangle \mapsto \langle \sigma', p + 1 \rangle} \text{STEPSEQ}$$

$$\frac{}{\langle \sigma, p \rangle \mapsto \langle \sigma, \perp \rangle} \text{CRASH}$$

$$\frac{\sigma' = \text{APPLY}(P[p], \sigma)}{\langle \sigma, p \rangle \mapsto \langle \sigma', p + 1 \rangle} \text{STEP}$$

$$\frac{}{\langle \sigma, p \rangle \mapsto \langle \sigma, \perp \rangle} \text{CRASH}$$

$$\frac{\sigma' = \text{PARTIALFLUSH}(\sigma)}{\langle \sigma, p \rangle \mapsto \langle \sigma', p \rangle} \text{NONDET}$$

**Figure:** An operational model for SCC

**Figure:** An operational model for ext4

# FERRITE


- ▶ Suite for reasoning about crash-consistency models
- ▶ Consists of two tools to explore all possible crash behaviors of a litmus test:
  - ▶ **Explicit enumerator:** executes litmus tests against actual file system implementations
  - ▶ **Bounded model checker:** executes litmus tests symbolically against an axiomatic specification



# FERRITE

Explicit Enumerator: Executing tests against file system implementations


1. Enumerator forwards litmus test on guest OS



# FERRITE

Explicit Enumerator: Executing tests against file system implementations


1. Enumerator forwards litmus test on guest OS
2. Guest OS executes litmus test



# FERRITE

Explicit Enumerator: Executing tests against file system implementations


1. Enumerator forwards litmus test on guest OS
2. Guest OS executes litmus test
3. File system issues events to the virtual disk



# FERRITE

## Explicit Enumerator: Executing tests against file system implementations


1. Enumerator forwards litmus test on guest OS
2. Guest OS executes litmus test
3. File system issues events to the virtual disk
4. FERRITE records events in a trace



# FERRITE

## Explicit Enumerator: Executing tests against file system implementations


1. Enumerator forwards litmus test on guest OS
2. Guest OS executes litmus test
3. File system issues events to the virtual disk
4. FERRITE records events in a trace
5. Enumerator prod. all possible reorderings and prefixes



# FERRITE

## Explicit Enumerator: Executing tests against file system implementations


1. Enumerator forwards litmus test on guest OS
2. Guest OS executes litmus test
3. File system issues events to the virtual disk
4. FERRITE records events in a trace
5. Enumerator prod. all possible reorderings and prefixes
6. Each prefix prod. a disk image file corresponding to a disk state after a crash



# FERRITE

## Explicit Enumerator: Executing tests against file system implementations

1. Enumerator forwards litmus test on guest OS
2. Guest OS executes litmus test
3. File system issues events to the virtual disk
4. FERRITE records events in a trace
5. Enumerator prod. all possible reorderings and prefixes
6. Each prefix prod. a disk image file corresponding to a disk state after a crash
7. Enumerator mounts disk image and checks predicates in the given litmus test



# FERRITE

Bounded model checker: Executing tests against specifications

- ▶ Input:
  - ▶ litmus test  $P$
  - ▶ axiomatic specification of a crash-consistency model  $M$
- ▶ Checker prod. set of crash traces
$$T = \{c_P \mid c_P \text{ is a crash trace of } t_P \text{ and } M(t_P, \tau_P)\}$$
- ▶  $\forall c_P \in T$ : checks whether  $c_P$  satisfies the predicates in  $P$



# Conclusion

The problem:

- ▶ POSIX standards under-specify guarantees that file systems should provide in the face of crashes

# Conclusion

The problem:

- ▶ POSIX standards under-specify guarantees that file systems should provide in the face of crashes
- ▶ Application writers can only guess how to achieve data integrity *and* improve performance

# Conclusion

The problem:

- ▶ POSIX standards under-specify guarantees that file systems should provide in the face of crashes
- ▶ Application writers can only guess how to achieve data integrity *and* improve performance

The solution:

- ▶ *Crash-consistency models* (i.e. litmus test and formal specifications) build a contract between the file system and the application

# Outlook

## Synthesis: Proof-of-concept

- ▶ Crash-consistency specifications allow to *synthesize* desired crash-safety properties

# Outlook

## Synthesis: Proof-of-concept

- ▶ Crash-consistency specifications allow to *synthesize* desired crash-safety properties
- ▶ Application writer develops  $P$  assuming SCC

# Outlook

## Synthesis: Proof-of-concept

- ▶ Crash-consistency specifications allow to *synthesize* desired crash-safety properties
- ▶ Application writer develops  $P$  assuming SCC
- ▶ Synthesizer transforms  $P$ , by inserting a minimal set of barriers (i.e. `fsync`)

# Outlook

## Synthesis: Proof-of-concept

- ▶ Crash-consistency specifications allow to *synthesize* desired crash-safety properties
- ▶ Application writer develops  $P$  assuming SCC
- ▶ Synthesizer transforms  $P$ , by inserting a minimal set of barriers (i.e. `fsync`)
- ▶ Resulting program  $P'$  behaves just like  $P$  under a given weaker crash-consistency model (e.g. `ext4`)